

## Gen AI Primer 2025

1. In a neural network, what does a neuron compute?

ANS: A weighted sum followed by an activation function

2. Which function introduces non-linearity in a neural network?

ANS: Activation Function

3. Which of the following is NOT a layer type in a typical neural network?

ANS: Quantum Layer

4. Which application of ML is used to detect unusual patterns in data?

ANS: Anomaly Detection

5. In which type of ML does an agent learn by interacting with an environment?

ANS: Reinforcement Learning

6. Which of the following is a common activation function in neural networks?

ANS: ReLU (Rectified Linear Unit)

7. How is a neural network's performance typically evaluated during training?

ANS: Using a validation set

8. Which of the following is a challenge in training deep neural networks?

ANS: Vanishing/Exploding gradients

9. What is the role of the loss function in training a neural network?

ANS: To quantify the difference between predicted and actual values

10. What is the primary purpose of backpropagation?

ANS: Adjusting weights based on the error

1. What is the primary purpose of a loss function in training neural networks?

ANS: To quantify the difference between predicted and actual values

2. Which of the following is a technique to prevent overfitting in neural networks?

ANS: Dropout

3. Which of the following is NOT a common machine learning algorithm?

ANS: Quantum Entanglement

4. What is the main difference between regression and classification?

ANS: Regression predicts a continuous output, Classification predicts a discrete label

5. What is the primary goal of machine learning?

ANS: To allow computers to learn from data

6. In the context of neural networks, what does the term "backpropagation" refer to?

ANS: The method of adjusting weights based on the error

7. Which component of a neural network is responsible for combining inputs and passing them to the next layer?

ANS: Neuron (or Node)

8. Which activation function outputs a value between 0 and 1?

ANS: Sigmoid

9. Which of the following is NOT a type of machine learning?

ANS: Recursive Learning

10. Which application of ML is used to group similar items?

ANS: Clustering

1. Which of the following is NOT a direct application of the Transformer architecture?

ANS: Image recognition

2. Which generative model introduced a stochastic layer that models data in a latent space?

ANS: VAE

3. What is the primary advantage of Transformers over RNNs in terms of processing sequences?

ANS: Parallel Processing

4. Which AI model series by OpenAI, based on the Transformer architecture, is known for generating highly coherent content?

ANS: GPT series

5. What mechanism allows the Transformer model to weigh the importance of different words in a sequence?

ANS: Self-Attention Mechanism

6. Which model can transform horse photos into zebra photos without direct comparison?

ANS: CycleGAN

7. What is the main innovation introduced by the "Attention Is All You Need" paper?

ANS: Transformer architecture

8. Which model is known for its rules for creating stable and effective AI image-makers?

ANS: DCGAN

9. Which model demonstrated that using larger architectures can produce better images?

ANS: BigGAN

10. In the context of GANs, what is the role of the Discriminator?

ANS: To distinguish between real and generated data

1. Who introduced Generative Adversarial Networks (GANs)?

ANS: Ian Goodfellow

2. What is the primary purpose of generative models?

ANS: Generating new data

3. Which of the following research papers is foundational for Variational Autoencoders (VAEs)?

ANS: "Generative Adversarial Nets"

4. Which model uses a probabilistic approach to encode and decode data?

ANS: VAE

5. Which of the following is NOT a direct application of GANs but rather an outcome of its influence?

ANS: Reinforcement learning in game playing

6. In which year were Generative Adversarial Networks (GANs) first introduced?

ANS: 2014

7. Which architecture is primarily associated with attention mechanisms?

ANS: Transformer

8. What are the two main components of a GAN?

ANS: Generator and Discriminator

9. Which model marked a significant milestone in the use of transformers in NLP?

ANS: BERT

10. Which pioneering research in Generative AI specifically emphasized the generation of text sequences?

ANS: "Sequence to Sequence Learning with Neural Networks"

1. How might AI transform the fitting experience in fashion retail?

ANS: By allowing customers to virtually "try on" clothes

2. What is a potential challenge of over-relying on AI in fashion?

ANS: Overshadowing human creativity and intuition

3. Which is a potential future application of Generative AI in fashion?

ANS: Hyper-personalized clothing

4. What could be a futuristic application of Generative AI in creating personalized clothing?

ANS: By designing clothing based on an individual's mood

5. Which brand has utilized AI for generating new clothing designs?

ANS: StitchFix

6. How can Generative AI contribute to sustainable fashion?

ANS: By designing sustainable fabrics

7. What aspect of AI in fashion raises concerns about user privacy?

ANS: Personalized shopping experiences

8. In the context of sustainability, how might AI be used in material design?

ANS: By simulating and designing eco-friendly materials

9. Who is a prominent virtual influencer mentioned in the case study?

ANS: Lil Miquela

10. How can Generative AI potentially impact inventory management in fashion?

ANS: By predicting fashion trends with higher accuracy

11. Beyond body image standards, what's another ethical concern regarding virtual models?

ANS: The authenticity of virtual influencers in advertising campaigns

12. Why might there be concerns about an over-reliance on data-driven AI in fashion?

ANS: Fashion inherently values human creativity and intuition

13. How does AI enhance the shopping experience?

ANS: By providing personalized product recommendations

14. What ethical concern arises from the use of virtual models in fashion?

ANS: Questions about body image standards and job losses for human models

15. What is the global valuation of the fashion industry?

ANS: \$2.5 trillion

1. If an AI system is designed to label images of cats and dogs, it is primarily a \_\_\_\_\_ model.

ANS: Discriminative

2. Which AI type is best for predicting outcomes?

ANS: Regression AI

3. Which is NOT a real-world application of Generative AI?

ANS: Predicting stock market prices

4. Which of the following is a direct application of Generative AI in the entertainment industry?

ANS: Creating realistic CGI characters

5. In the context of AI, which model type is more concerned with the underlying distribution of data?

ANS: Generative AI

6. What is Generative AI primarily used for?

ANS: Generating new data

7. Generative AI can be used to create which of the following?

ANS: New artworks and music pieces

8. How does Generative AI differ from Classification AI?

ANS: It generates new data rather than categorizing existing data

9. Which statement best describes the role of Generative AI?

ANS: It focuses on generating data based on learned patterns

10. Why is Generative AI considered significant in the realm of artificial intelligence?

ANS: It can produce new, previously unseen data samples

1. Which of the following fields can utilize Generative AI to create new, original content or simulations?

ANS: Art and Music

2. Which AI type primarily focuses on labeling data?

ANS: Supervised AI

3. Generative AI is closely related to which type of models?

ANS: Generative models

4. Which type of AI is primarily concerned with how data is generated rather than how it's separated?

ANS: Generative AI

5. In which application is Generative AI NOT typically used?

ANS: Automating customer service chats

6. What does AI stand for?

ANS: Artificial Intelligence

7. Which of the following is NOT a type of AI?

ANS: Generative Art

8. Which statement best defines Generative AI?

ANS: Generative AI that can generate new data samples

1. Which of the following is NOT a property of likelihood?

ANS: It is normalized like a probability

2. Which model type aims to capture the joint probability  $P(x, y)$ ?

ANS: Generative Model

3. Which claim regarding generative models isn't true?

ANS: They always require labeled data for training

4. Which of the following best describes the difference between generative and discriminative models?

ANS: Generative models learn the data distribution, while discriminative models learn the decision boundary

5. Within the architecture of Generative Adversarial Networks (GANs), which duo of fundamental elements are paramount?

ANS: Generator and Discriminator

6. For what tasks can generative models be applied?

ANS: Data generation, denoising, inpainting, and more

7. What does a probability distribution provide?

ANS: A mathematical description of outcomes for a random variable

8. What's a significant hurdle when training GANs?

ANS: Mode collapse

9. Within generative models, what function does the discriminator serve in GANs?

ANS: To distinguish between real and generated data

10. How is the likelihood of data given a model symbolized?

ANS:  $P(\text{data} \mid \text{model})$

1. If a model is better at distinguishing between classes rather than generating data, it is likely a \_\_\_\_\_.

ANS: Discriminative model

2. Which model type is primarily concerned with determining  $P(y \mid x)$ ?

ANS: Discriminative Model

3. What is the primary goal of generative models in AI?

ANS: To generate new data samples

4. Which statement best differentiates generative from discriminative models?

ANS: Generative models learn the joint probability distribution, while discriminative models learn the conditional probability

5. Which of the following is crucial for understanding the behavior of generative models?

ANS: Probability distributions and likelihood

6. In the context of generative models, what does  $P(x)$  represent?

ANS: The probability distribution of the data  $x$

7. Which of the following is NOT a generative model?

ANS: Support Vector Machines

8. Generative models are primarily used for which of the following tasks?

ANS: Generating new data samples similar to the input data

9. What does likelihood measure in the context of a model?

ANS: How well the model explains the observed data

10. In the context of models, what does  $P(x \mid y)$  typically represent?

ANS: The probability of  $x$  given  $y$

1. What is a primary application of VAEs mentioned in the case study?

ANS: Anomaly Detection

2. What is the y-axis label of the chart visualizing the error?

ANS: Reconstruction Error

3. Which is NOT a challenge in implementing VAEs for this use-case?

ANS: Increasing data storage costs

4. For how many epochs is the VAE trained?

ANS: 50

5. What criterion is used to determine if a data point is anomalous?

ANS: If its error is above the 99th percentile

6. What is the VAE trained to learn effectively?

ANS: A compressed representation of the data

7. Why is understanding the VAE's outputs challenging?

ANS: They can be complex and non-intuitive

8. In the VAE, what does the sampling function introduce?

ANS: Randomness

9. What two components combine to form the VAE's loss?

ANS: MSE and KL divergence

10. Which of the following is NOT an attribute in the given data?

ANS: Humidity

11. What type of dataset does the manufacturing plant collect?

ANS: Time Series Dataset

12. How is the data divided for training the VAE?

ANS: 80-20

13. What does the VAE attempt to minimize during training?

ANS: Loss

14. Why is data preprocessing required before training the VAE?

ANS: To ensure it is suitable for training

15. Over time, due to certain changes, what might be required of the VAE model?

ANS: Continuous adaptation

1. In which application can VAEs detect unusual patterns?

ANS: Anomaly detection in industrial equipment

2. Which of the following is a key component of the VAE loss function?

ANS: KL divergence

3. In which application might VAEs be used to enhance image quality?

ANS: Medical imaging

4. Why is the reparameterization trick important in VAEs?

ANS: It allows backpropagation through random nodes

5. Which optimization technique is commonly used with VAEs?

ANS: Stochastic gradient descent (SGD)

6. Autoencoders primarily focus on which aspect of data?

ANS: Reconstruction

7. Which of the following is NOT a typical use case for VAEs?

ANS: Real-time speech translation

8. Why is variational inference used in VAEs?

ANS: To approximate intractable posterior distributions

9. How do VAEs differ from traditional autoencoders?

ANS: VAEs introduce randomness via a probabilistic layer

10. What do VAEs use to generate a distribution over latent variables?

ANS: Variational inference

1. Why are autoencoders considered generative models?

ANS: They can reconstruct and generate data similar to the input

2. Reparameterization trick is used to...

ANS: Deal with the non-differentiability of sampling in VAEs

3. What does VAE stand for?

ANS: Variational Autoencoder

4. What is the primary role of autoencoders in generative modeling?

ANS: Data compression and reconstruction

5. Which application does NOT typically use VAEs?

ANS: Text summarization

6. Why is the reparameterization trick crucial in training VAEs?

ANS: It allows backpropagation through stochastic nodes

7. In the context of Variational Autoencoders (VAEs), what does variational inference help achieve?

ANS: Approximation of complex posterior distributions

8. Which component of the VAE loss function ensures the latent variables adhere to a standard distribution?

ANS: KL divergence

9. In which application might you use a VAE for generating new, coherent samples?

ANS: Designing virtual fashion items

10. Which of the following is NOT a type of autoencoder?

ANS: Supervised autoencoder